

Kubernetes Operations Guide

Department: Engineering

Cluster Setup

Production workloads run on AWS EKS across 3 availability zones. Each service is deployed as a separate Kubernetes Deployment with a minimum of 2 replicas.

Autoscaling

Horizontal Pod Autoscaling is configured to trigger when CPU utilization exceeds 70 percent for more than 5 minutes, scaling up to a maximum of 10 replicas.

Resource Limits

Every pod must define CPU and memory requests and limits in its manifest. Pods without defined limits will be rejected by the admission controller.